

HisabDo AI Copilot

Identify Required Website/Content Data

Prepared by: Laiba Yousuf

Track/Department: AI/ML-3

1. Scope and Approach

HisabDo does not have a standalone marketing website — it is distributed as a mobile app (Google Play, package com.usman.hisabdo). "Website content" for this task is therefore interpreted as the two real public/internal sources available: (1) the live Google Play Store listing, which is the closest thing to official public-facing content, and (2) the internal AI knowledge base already authored inside this repository (HisabDo_RAG_Chatbot/knowledge-base/), which contains 14 markdown files covering FAQs, features, guides, and blog content, built and evaluated by a teammate as part of the RAG pipeline (Hit@3: 100% on a 12-query test set).

2. Content Sources Reviewed

- Google Play Store listing for HisabDo: Khata & Ledger (live, public) — app description, feature bullets, category positioning
- knowledge-base/faqs/ — 4 files: general.md, accounts.md, transactions.md, backup.md
- knowledge-base/features/ — 5 files: khata.md, customers.md, analytics.md, pdf-export.md, voice-entry.md
- knowledge-base/guides/ — 3 files: getting-started.md, expense-tracking.md, managing-customers.md
- knowledge-base/blogs/ — 2 files: daily-khata.md, digital-ledger.md
- RAG_Knowledge_Base_Spec.docx — confirms this 14-file structure matches Section 10 of the team's architecture document

3. Product Descriptions (collected)

The Play Store listing describes HisabDo as a khata app for managing udhar, ledger accounts, receivables, payables, expenses, and customer transactions, positioned for shopkeepers, freelancers, wholesalers, and home-based sellers. It highlights an offline-first, local-database architecture (no cloud dependency), biometric/PIN security, and positions itself against competitors on data privacy grounds. This description is usable as-is for a "What is HisabDo" knowledge-base entry, and is consistent with (not contradicted by) faqs/general.md, which already covers this question internally.

4. Feature Explanations (collected)

- Khata / ledger tracking — running customer balance from credit-given and payment-received entries (features/khata.md)

- Customer management — add customer, per-customer profile with balance and history (features/customers.md)
- Analytics — receivables total, monthly transaction totals, trend charts (features/analytics.md)
- PDF export — per-customer statement export/share (features/pdf-export.md)
- Voice entry — speech-to-text transaction pre-fill (features/voice-entry.md)
- From the Play listing, not yet in the internal KB: built-in calculator, WhatsApp/SMS payment reminders, multi-currency support (PKR/USD/INR), multi-language support (Urdu/English/Hindi/Arabic/Roman Urdu)

5. Service Information (collected)

- Backup & restore — automatic cloud sync when connected, manual "Backup Now" option (faqs/backup.md)
- Account creation — phone number or email sign-up (faqs/accounts.md)
- Multi-language and multi-currency support — confirmed in the Play listing; not yet reflected in the internal knowledge base

6. Help and Support Content (identified)

The internal FAQ set (4 files) already covers account setup, transactions, and backup. The Play Store listing references in-app "Help & Support," but that in-app content itself has not been extracted into the knowledge base and is listed as missing below.

7. Blog Articles (identified)

Two blog articles exist: "Why a Daily Khata Habit Matters for Small Businesses" and "From Paper to Digital: What Changes When You Switch Your Ledger." Both are evergreen educational content (not time-sensitive), suitable for direct inclusion once reviewed for the flagged language noted in Section 9.

8. Missing Documentation (identified)

- Pricing/plan tiers — Play listing mentions a free tier but full plan details are not documented anywhere in the repo
- Privacy policy / data handling detail — beyond the general "offline-first, no cloud risk" marketing claim
- Terms of service
- Troubleshooting / error-message guide — no dedicated content exists for common failure cases
- In-app Help & Support article content — referenced by the Play listing but not extracted
- Onboarding/tutorial content beyond guides/getting-started.md (e.g. for voice entry or PDF export specifically)

9. Outdated or Unclear Content (flagged)


Three passages use conditional or uncertain language, which is a specific risk for AI ingestion: a knowledge base entry phrased as "check whether your app has this feature" will cause the AI assistant to

pass that same uncertainty on to a real user instead of giving a confident answer. These should be verified against the current live app before inclusion.

File	Flagged text	Why it needs review
guides/expense-tracking.md	"check whether your app version has a dedicated Expenses section"	Author was unsure if the feature exists in the current app — needs verification against the live app before ingestion.
faqs/accounts.md	"check whether your plan supports team/staff accounts"	Same pattern — conditional/uncertain language about plan-specific behavior.
blogs/daily-khata.md	"a shopkeeper might undercount out of uncertainty"	Speculative phrasing (fine for a blog's narrative tone, but should not be treated as a factual product claim if reused elsewhere).

10. Content to Include in the AI Knowledge Base

Recommended for inclusion as-is (already reviewed, already tested in the working RAG pipeline with 100% Hit@3): all 4 FAQ files, all 5 feature files, and both blog articles. Recommended for inclusion after the fixes noted in Section 9: getting-started.md, expense-tracking.md, managing-customers.md, accounts.md. Recommended for new collection before inclusion: Play Store listing content (calculator, reminders, multi-currency, multi-language — not yet in any file), plus the missing categories in Section 8.



11. Content Inventory and Categorization

Mapped against the task's suggested knowledge sources (FAQs, Website Pages, Feature Descriptions, Help Content, Product Guides, Documentation, Blog Articles, Support Information, User Guides):

Category	Source(s) Identified	Status	Priority
FAQs	faqs/general.md, accounts.md, transactions.md, backup.md	Exists (internal KB)	High
Feature Descriptions	features/khata.md, customers.md, analytics.md, pdf-export.md, voice-entry.md	Exists (internal KB)	Very High
Product Guides / User Guides	guides/getting-started.md, expense-tracking.md, managing-customers.md	Exists (internal KB) — 1 flagged for review	High

Category	Source(s) Identified	Status	Priority
Blog Articles	blogs/daily-khata.md, digital-ledger.md	Exists (internal KB)	Medium
Website Pages / Product Description	Google Play Store listing (public, live)	Exists externally — not yet imported	Very High
Support Information	In-app "Help & Support" (referenced in Play listing, content not extracted)	Missing — needs collection	High
Documentation (technical/policy)	Pricing/plans detail, privacy policy, terms of service	Missing	Medium
Troubleshooting Content	Error messages, common-issue resolutions	Missing	Medium
Customer Management	features/customers.md	Verified	Very High

12. Preparing Content for Cleaning and Chunking


The existing pipeline (document_loader.py, chunker.py) already cleans markdown syntax and chunks along "###" section headers, producing one chunk per question/topic unit. Any newly collected content (Play Store feature list, pricing, troubleshooting) should be written in the same markdown format with "###" subheadings per topic, so it is immediately compatible with the existing chunker without changes to the pipeline code.

document_loader.py

chunker.py

embedder.py

RAG Knowledge Base Ingestion Pipeline



13. Verified vs. Unverified Content

Verified: the Play Store listing (live, official, current as of this review) and the 14 internal knowledge-base files (authored by the RAG/KB engineer, evaluated with a real 12-query test set). Unverified: the three flagged passages in Section 9, and any assumption about in-app Help & Support content, plan/pricing tiers, or troubleshooting steps, since none of these have a documented source yet.

14. Recommended Content Update Process

- Add a source, owner, and last-reviewed date to each knowledge-base file
- When a file is edited, re-run chunker.py and embedder.py so the vector index reflects the change
- Periodically re-check the Play Store listing for description/feature changes, since it is the only live external source
- Route any new pricing, policy, or legal content through a designated content owner before ingestion, given its higher accuracy risk

15. Summary — Expected Output

This document is the complete list of website and product content required for AI knowledge integration: 14 existing, largely reusable knowledge-base files; one external source (the Play Store listing) with content not yet imported; three specific passages flagged for accuracy review; and six missing content categories (pricing, privacy policy, terms of service, troubleshooting, in-app help content, and expanded tutorials) that should be collected before the knowledge base is considered complete.
